

exit(3) — Linux manual page

[NAME](#) | [SYNOPSIS](#) | [DESCRIPTION](#) | [RETURN VALUE](#) | [ATTRIBUTES](#) | [CONFORMING TO](#) | [NOTES](#) | [SEE ALSO](#) | [COLOPHON](#)

 EXIT(3)**Linux Programmer's Manual****EXIT(3)****NAME** [top](#)

exit - cause normal process termination

SYNOPSIS [top](#)

```
#include <stdlib.h>
```

```
noreturn void exit(int status);
```

DESCRIPTION [top](#)

The **exit()** function causes normal process termination and the least significant byte of *status* (i.e., *status* & 0xFF) is returned to the parent (see [wait\(2\)](#)).

All functions registered with [atexit\(3\)](#) and [on_exit\(3\)](#) are called, in the reverse order of their registration. (It is possible for one of these functions to use [atexit\(3\)](#) or [on_exit\(3\)](#) to register an additional function to be executed during exit processing; the new registration is added to the front of the list of functions that remain to be called.) If one of these functions does not return (e.g., it calls [_exit\(2\)](#), or kills itself with a signal), then none of the remaining functions is called, and further exit processing (in particular, flushing of [stdio\(3\)](#) streams) is abandoned. If a function has been registered multiple times using [atexit\(3\)](#) or [on_exit\(3\)](#), then it is called as many times as it was registered.

All open [stdio\(3\)](#) streams are flushed and closed. Files created by [tmpfile\(3\)](#) are removed.

The C standard specifies two constants, **EXIT_SUCCESS** and **EXIT_FAILURE**, that may be passed to **exit()** to indicate successful or unsuccessful termination, respectively.

RETURN VALUE [top](#)

The **exit()** function does not return.

ATTRIBUTES [top](#)

For an explanation of the terms used in this section, see [attributes\(7\)](#).

Interface	Attribute	Value
exit()	Thread safety	MT-Unsafe race:exit

The **exit()** function uses a global variable that is not protected, so it is not thread-safe.

CONFORMING TO [top](#)

POSIX.1-2001, POSIX.1-2008, C89, C99, SVr4, 4.3BSD.

NOTES [top](#)

The behavior is undefined if one of the functions registered using [atexit\(3\)](#) and [on_exit\(3\)](#) calls either **exit()** or [longjmp\(3\)](#). Note that a call to [execve\(2\)](#) removes registrations created using [atexit\(3\)](#) and [on_exit\(3\)](#).

The use of **EXIT_SUCCESS** and **EXIT_FAILURE** is slightly more portable (to non-UNIX environments) than the use of 0 and some nonzero value like 1 or -1. In particular, VMS uses a different convention.

BSD has attempted to standardize exit codes (which some C libraries such as the GNU C library have also adopted); see the file [<sysexits.h>](#).

After **exit()**, the exit status must be transmitted to the parent process. There are three cases:

- If the parent has set **SA_NOCLDWAIT**, or has set the **SIGCHLD** handler to **SIG_IGN**, the status is discarded and the child dies immediately.
- If the parent was waiting on the child, it is notified of the exit status and the child dies immediately.
- Otherwise, the child becomes a "zombie" process: most of the process resources are recycled, but a slot containing minimal information about the child process (termination status, resource usage statistics) is retained in process table. This allows the parent to subsequently use [waitpid\(2\)](#) (or similar) to learn the termination status of the child; at that point the zombie process slot is released.

If the implementation supports the **SIGCHLD** signal, this signal is sent to the parent. If the parent has set **SA_NOCLDWAIT**, it is undefined whether a **SIGCHLD** signal is sent.

Signals sent to other processes

If the exiting process is a session leader and its controlling terminal is the controlling terminal of the session, then each process in the foreground process group of this controlling terminal is sent a **SIGHUP** signal, and the terminal is disassociated from this session, allowing it to be acquired by a new controlling process.

If the exit of the process causes a process group to become orphaned, and if any member of the newly orphaned process group is stopped, then a **SIGHUP** signal followed by a **SIGCONT** signal will be sent to each process in this process group. See [setpgid\(2\)](#) for an explanation of orphaned process groups.

Except in the above cases, where the signalled processes may be children of the terminating process, termination of a process does *not* in general cause a signal to be sent to children of that process. However, a process can use the [prctl\(2\)](#) **PR_SET_PDEATHSIG** operation to arrange that it receives a signal if its parent terminates.

SEE ALSO [top](#)

[_exit\(2\)](#), [get_robust_list\(2\)](#), [setpgid\(2\)](#), [wait\(2\)](#), [atexit\(3\)](#), [on_exit\(3\)](#), [tmpfile\(3\)](#)

COLOPHON [top](#)

This page is part of release 5.13 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.

Linux

2021-03-22

EXIT(3)

Pages that refer to this page: [man\(1\)](#), [_exit\(2\)](#), [kill\(2\)](#), [vfork\(2\)](#), [wait\(2\)](#), [abort\(3\)](#), [assert\(3\)](#), [assert_perror\(3\)](#), [atexit\(3\)](#), [err\(3\)](#), [error\(3\)](#), [on_exit\(3\)](#), [pthread_create\(3\)](#), [pthread_detach\(3\)](#), [pthread_exit\(3\)](#), [sd_bus_set_exit_on_disconnect\(3\)](#), [setjmp\(3\)](#), [stdin\(3\)](#), [stdio\(3\)](#), [tmpfile\(3\)](#)

[Copyright and license for this manual page](#)

HTML rendering created 2021-08-27 by [Michael Kerrisk](#), author of *The Linux Programming Interface*, maintainer of the [Linux man-pages project](#).

For details of in-depth **Linux/UNIX system programming training courses** that I teach, look [here](#).

Hosting by [jambit GmbH](#).

